

An Experimentation Toolkit for Robotics Control and Manipulation Tasks using Reinforcement Learning Algorithms

A Robot Learning Gym

Ashwin Reddy¹

¹The Harker School

March 22, 2017

Introduction

Goals of the project:


- ▶ Build a toolkit for Deep Robot RL experimentation using a simulator
 - ▶ Collect models and tasks online
 - ▶ Create an API to interface them
- ▶ Create a simple benchmark

Motivation

- ▶ Such a tool does not currently exist, but it would be helpful
- ▶ Standard metrics missing for robot learning algorithms
- ▶ Enables more Sim-to-Real learning
- ▶ Does not require a real robot to operate

Background: RL


- ▶ Traditional Reinforcement Learning Paradigm (MDPs):



- ▶ How can we train the agent to pick actions to maximize reward?
- ▶ Goal: Learn parameters θ to stochastic policy $\pi_{\theta}(a_t|s_t)$
- ▶ Objective η is $\max \mathbb{E} \left[\gamma^t \int_t^H r(t) dt \right]$

Background: ML

- ▶ Machine Learning Paradigm (using Neural Networks):



- ▶ Goal: make predictions
- ▶ Given many samples of $(x, f(x))$, learn $h_{\theta}(x) \approx f(x)$

Current Work

- ▶ Algorithms are starting to use deep machine learning:
 - ▶ Google DeepMind's Deep-Q-Learning which has learned to play Atari games
 - ▶ UC Berkeley's Guided Policy Search which taught a humanoid robot to assemble toy airplanes
- ▶ Different ML Architectures:
 - ▶ Modular Neural Networks
 - ▶ Progressive Neural Networks
 - ▶ Recurrent Neural Networks

What makes Deep RL for Robotics hard?

- ▶ Lack of large datasets (i.e: poor supervision)
- ▶ Noisy input and can't be sure of outputs
- ▶ Occlusions
- ▶ Running on a real robot is expensive
- ▶ Curse of dimensionality
- ▶ Sparse and time-delayed rewards
- ▶ Exploration vs. Exploitation
- ▶ Continuous control harder than discretized actions

Current Tools

- ▶ OpenAI Gym: Platform for testing general RL algorithms
- ▶ MuJoCo: Efficient simulator used by OpenAI Gym
- ▶ OpenAI RL Lab: implements some common RL algorithms
- ▶ Guided Policy Search: A type of Deep RL algorithm that has a package available online, also uses MuJoCo
- ▶ DeepMind's Lab: 3D first-person games for testing Deep RL algorithms

All have elements of robot learning, but don't provide a framework for it


Experimental Design

- ▶ Collect MuJoCo Robot Models
- ▶ Build an open-source framework with models and other tools
 - ▶ OpenAI MuJoCo Python bindings and Gym
- ▶ Run RL algorithms

Data Collection

- ▶ Task: Peg Insertion with $r = \frac{1}{||x-x^*||}$
- ▶ Methods
 - ▶ Random: take a sample from the action space
 - ▶ Cross-Entropy Method: sample parameters, evaluate, and reuse best ones
 - ▶ Policy Gradient Method: Use trajectory $\langle s, a, r, s' \rangle$ to $\theta_{k+1} \leftarrow \theta_k + \alpha \nabla_{\theta} \eta$
 - ▶ Guided Policy Search: uses trajectory optimization to train the neural net based policy

Graphs (Reward vs Time)



- ▶ Collected using TensorFlow's TensorBoard Visualizer, using a 0.6 smooth
- ▶ Teal line is policy gradient, orange line is cross-entropy, purple is random
- ▶ Random has erratic behavior and inconsistent results as expected
- ▶ Cross entropy is somewhat unstable but leads to more consistent results
- ▶ Policy Gradient seems to have reached a local optimum

Discussion and Future Work

- ▶ Continue to maintain the framework and create documentation
- ▶ Experiment with modular networks
- ▶ Benchmark more deep learning based algorithms
- ▶ Try different tasks and rewards

Acknowledgements and References

- ▶ Thanks to
 - ▶ Mr. Martin Baynes for sponsoring this project
 - ▶ Soroush Nasiriany, an undergraduate researcher at UC Berkeley, for feedback on this project